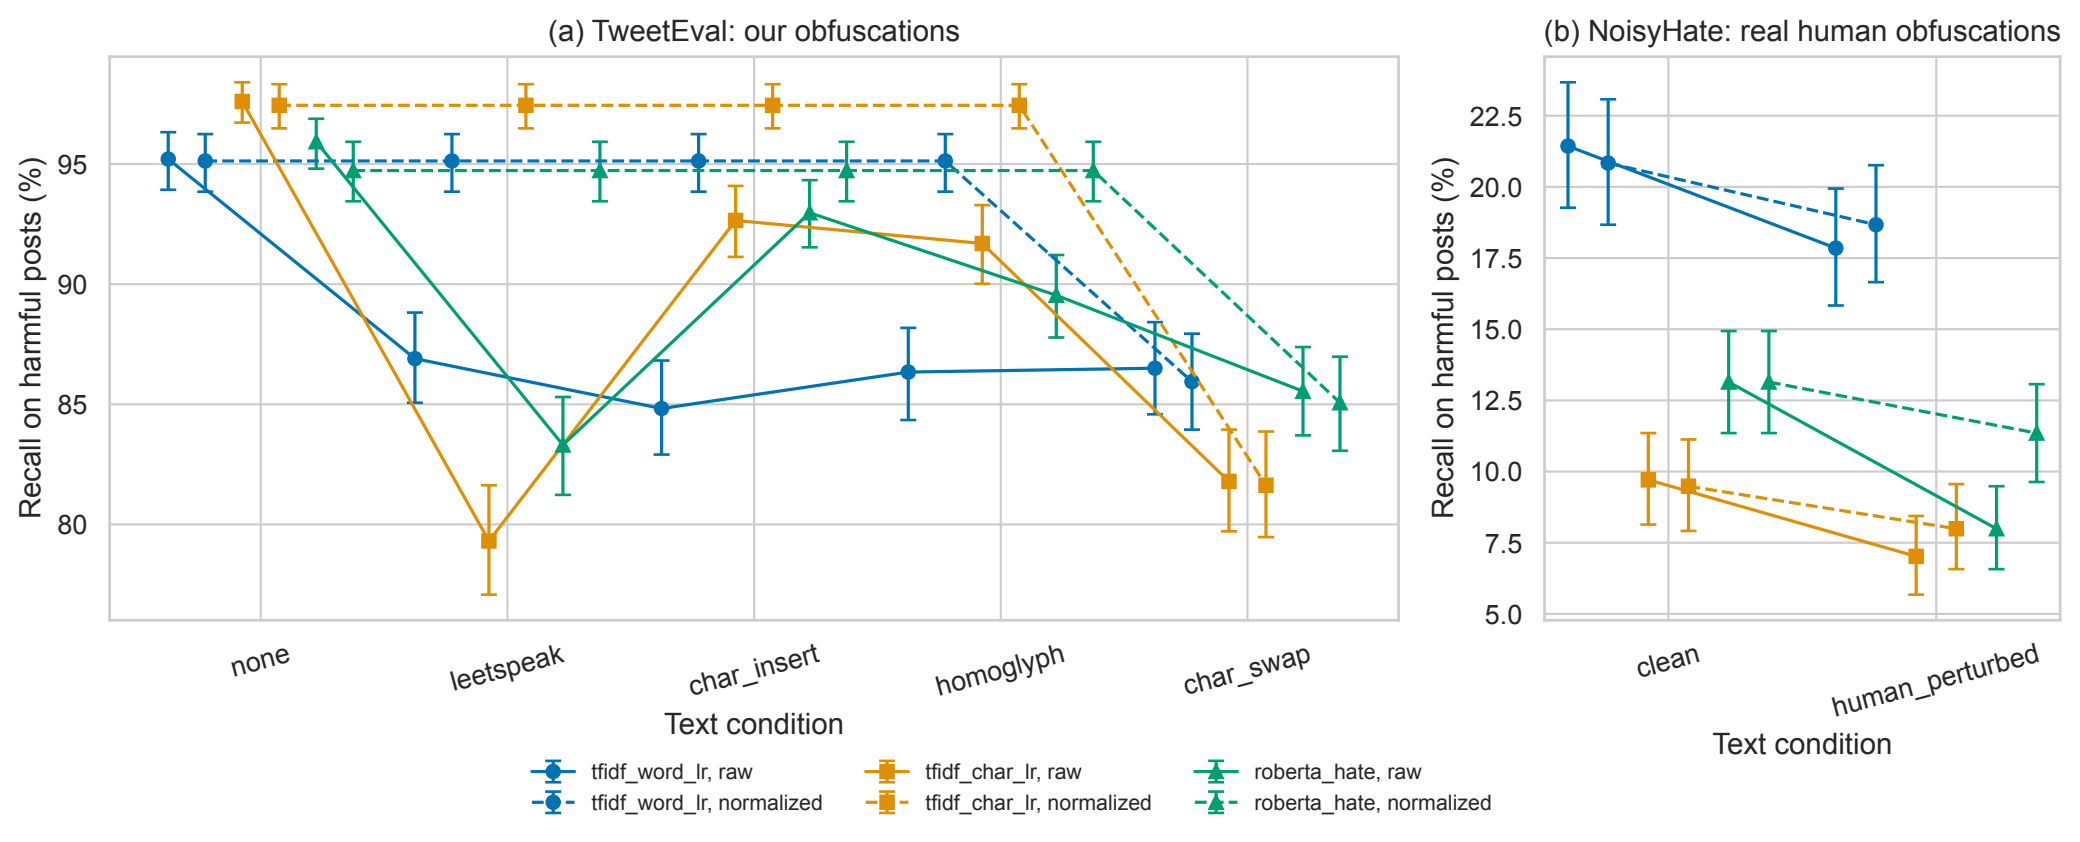
(a) TweetEval: our obfuscations
Recall on harmful posts (%)
95
90
85
80
none
leetspeak
char_insert
homoglyph
char_swap
Text condition
(b) NoisyHate: real human obfuscations
Recall on harmful posts (%)
22.5
20.0
17.5
15.0
12.5
10.0
7.5
5.0
clean
human_perturbed
Text condition
tfidf_word_lr, raw
tfidf_word_lr, normalized
tfidf_char_lr, raw
tfidf_char_lr, normalized
roberta_hate, raw
roberta_hate, normalized